

pandas_profiling_report	2
data_preparation	9
baseline_models	16
selected_features_models	25

Overview

Overview

Alerts 18

Reproduction

Dataset statistics

Number of variables	19
Number of observations	8760
Missing cells	0
Missing cells (%)	0.0%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	1.3 MiB
Average record size in memory	152.0 B

Variable types

Numeric	12
Categorical	7

Variables

Select Columns


Rented Bike Count

Real number (R)

HIGH CORRELATIONZEROS

Distinct	2166
Distinct (%)	24.7%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	704.60205

Minimum	0
Maximum	3556
Zeros	295
Zeros (%)	3.4%
Negative	0
Negative (%)	0.0%
Memory size	68.6 KiB




More details

Hour

Real number (R)

Distinct	24
Distinct (%)	0.3%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	11.5

Minimum	0
Maximum	23
Zeros	365
Zeros (%)	4.2%
Negative	0
Negative (%)	0.0%
Memory size	68.6 KiB




More details

Temperature(°C)

Real number (R)

Distinct	546
Distinct (%)	6.2%
Missing	0
Missing (%)	0.0%
Infinite	0

Minimum	-17.8
Maximum	39.4
Zeros	21
Zeros (%)	0.2%
Negative	1433




Temperature(°C)

Real number (ℝ)

Distinct	546
Distinct (%)	6.2%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	12.882922

Minimum	-17.8
Maximum	39.4
Zeros	21
Zeros (%)	0.2%
Negative	1433
Negative (%)	16.4%
Memory size	68.6 KiB


[More details](#)

Humidity(%)

Real number (ℝ)

Distinct	90
Distinct (%)	1.0%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	58.226256

Minimum	0
Maximum	98
Zeros	17
Zeros (%)	0.2%
Negative	0
Negative (%)	0.0%
Memory size	68.6 KiB


[More details](#)

Wind speed (m/s)

Real number (ℝ)

Distinct	65
Distinct (%)	0.7%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	1.7249087

Minimum	0
Maximum	7.4
Zeros	74
Zeros (%)	0.8%
Negative	0
Negative (%)	0.0%
Memory size	68.6 KiB


[More details](#)

Visibility (10m)

Real number (ℝ)

Distinct	1789
Distinct (%)	20.4%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	1436.8258

Minimum	27
Maximum	2000
Zeros	0
Zeros (%)	0.0%
Negative	0
Negative (%)	0.0%
Memory size	68.6 KiB

[More details](#)

Dew point temperature(°C)

Real number (ℝ)

Distinct	556
----------	-----


Minimum	-30.6
---------	-------

Dew point temperature(°C)

Real number (R)

Distinct	556
Distinct (%)	6.3%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	4.0738128

Minimum	-30.6
Maximum	27.2
Zeros	60
Zeros (%)	0.7%
Negative	3138
Negative (%)	35.8%
Memory size	68.6 KiB




More details

Solar Radiation (MJ/m2)

Real number (R)

Distinct	345
Distinct (%)	3.9%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	0.56911073

Minimum	0
Maximum	3.52
Zeros	4300
Zeros (%)	49.1%
Negative	0
Negative (%)	0.0%
Memory size	68.6 KiB



More details

Rainfall(mm)

Real number (R)

Distinct	61
Distinct (%)	0.7%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	0.14868721

Minimum	0
Maximum	35
Zeros	8232
Zeros (%)	94.0%
Negative	0
Negative (%)	0.0%
Memory size	68.6 KiB



More details

Snowfall (cm)

Real number (R)

Distinct	51
Distinct (%)	0.6%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	0.075068493

Minimum	0
Maximum	8.8
Zeros	8317
Zeros (%)	94.9%
Negative	0
Negative (%)	0.0%
Memory size	68.6 KiB



More details

Holiday

Categorical

Distinct	2	0	8328
		1	432

Holiday

Categorical

Distinct	2
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	68.6 KiB


[More details](#)

Functioning Day

Categorical

Distinct	2
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	68.6 KiB



[More details](#)

Day

Real number (ℝ)

Distinct	31
Distinct (%)	0.4%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	15.720548

Minimum	1
Maximum	31
Zeros	0
Zeros (%)	0.0%
Negative	0
Negative (%)	0.0%
Memory size	68.6 KiB



[More details](#)

Month

Real number (ℝ)

Distinct	12
Distinct (%)	0.1%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	6.5260274

Minimum	1
Maximum	12
Zeros	0
Zeros (%)	0.0%
Negative	0
Negative (%)	0.0%
Memory size	68.6 KiB


[More details](#)

Year

Categorical

HIGH CORRELATION **IMBALANCE**

Distinct	2
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%



Year

Categorical

HIGH CORRELATIONIMBALANCE

Distinct	2
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	68.6 KiB



More details

Autumn

Categorical

Distinct	2
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	68.6 KiB



More details

Spring

Categorical

Distinct	2
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	68.6 KiB



More details

Summer

Categorical

Distinct	2
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	68.6 KiB



More details

Winter

Categorical


Distinct	2
Distinct (%)	< 0.1%
Missing	0
Missing (%)	0.0%
Memory size	68.6 KiB



More details

Interactions

Rented Bike Count	Month
Hour	Rented Bike Count
Temperature(°C)	Hour
Humidity(%)	Temperature(°C)
Wind speed (m/s)	Humidity(%)
Visibility (10m)	Wind speed (m/s)
Dew point temperature(°C)	Visibility (10m)
Solar Radiation (MJ/m2)	Dew point temperature(°C)
Rainfall(mm)	Solar Radiation (MJ/m2)
Snowfall (cm)	Rainfall(mm)
Day	Snowfall (cm)
Month	Day




Correlations


Auto

Heatmap


Table



Missing values



Missing values



Sample

First rowsLast rows

	Rented Bike Count	Hour	Temperature(°C)	Humidity(%)	Wind speed (m/s)	Visibility (10m)	Dew point temperature(°C)
0	254	0	-5.2	37	2.2	2000	-17.6
1	204	1	-5.5	38	0.8	2000	-17.6
2	173	2	-6.0	39	1.0	2000	-17.7
3	107	3	-6.2	40	0.9	2000	-17.6
4	78	4	-6.0	36	2.3	2000	-18.6
5	100	5	-6.4	37	1.5	2000	-18.7
6	181	6	-6.6	35	1.3	2000	-19.5
7	460	7	-7.4	38	0.9	2000	-19.3
8	930	8	-7.6	37	1.1	2000	-19.8
9	490	9	-6.5	27	0.5	1928	-22.4

data_preparation

March 13, 2023

Data Preparation

For the SeoulBikeData.csv dataset, I will convert the categorical attributes to quantitative attributes such that they are in a valid format for use in building the regression models.

```
[46]: # Import the required modules
import pandas as pd
```

```
[47]: # Read in the dataset into a pandas dataframe

seoul_bike_data = pd.read_csv("SeoulBikeData.csv", encoding="ansi")
seoul_bike_data.head()
```

```
[47]:
```

	Date	Rented Bike Count	Hour	Temperature(°C)	Humidity(%)	\
0	01/12/2017	254	0	-5.2	37	
1	01/12/2017	204	1	-5.5	38	
2	01/12/2017	173	2	-6.0	39	
3	01/12/2017	107	3	-6.2	40	
4	01/12/2017	78	4	-6.0	36	

	Wind speed (m/s)	Visibility (10m)	Dew point temperature(°C)	\
0	2.2	2000	-17.6	
1	0.8	2000	-17.6	
2	1.0	2000	-17.7	
3	0.9	2000	-17.6	
4	2.3	2000	-18.6	

	Solar Radiation (MJ/m2)	Rainfall(mm)	Snowfall (cm)	Seasons	Holiday	\
0	0.0	0.0	0.0	Winter	No Holiday	
1	0.0	0.0	0.0	Winter	No Holiday	
2	0.0	0.0	0.0	Winter	No Holiday	
3	0.0	0.0	0.0	Winter	No Holiday	
4	0.0	0.0	0.0	Winter	No Holiday	

	Functioning Day
0	Yes
1	Yes
2	Yes
3	Yes

4 Yes

For this dataset, the categorical variables which I will transform are Date, Seasons, Holiday, and Functioning Day.

```
[48]: seoul_bike_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8760 entries, 0 to 8759
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Date                                8760 non-null   object
1   Rented Bike Count                  8760 non-null   int64
2   Hour                              8760 non-null   int64
3   Temperature(°C)                   8760 non-null   float64
4   Humidity(%)                       8760 non-null   int64
5   Wind speed (m/s)                  8760 non-null   float64
6   Visibility (10m)                  8760 non-null   int64
7   Dew point temperature(°C)         8760 non-null   float64
8   Solar Radiation (MJ/m2)           8760 non-null   float64
9   Rainfall(mm)                      8760 non-null   float64
10  Snowfall (cm)                     8760 non-null   float64
11  Seasons                           8760 non-null   object
12  Holiday                           8760 non-null   object
13  Functioning Day                    8760 non-null   object
dtypes: float64(6), int64(4), object(4)
memory usage: 958.2+ KB
```

1. Date Attribute First, for the Date attribute, I will transform it to get the Day, Month, and Year as these 3 new attributes will be quantitative.

```
[49]: # Use pandas to convert the Date attribute to a datetime data type
seoul_bike_data["Date"] = pd.to_datetime(seoul_bike_data["Date"], format="%d/%m/%Y")
seoul_bike_data.head()
```

```
[49]:
```

	Date	Rented Bike Count	Hour	Temperature(°C)	Humidity(%)	\
0	2017-12-01	254	0	-5.2	37	
1	2017-12-01	204	1	-5.5	38	
2	2017-12-01	173	2	-6.0	39	
3	2017-12-01	107	3	-6.2	40	
4	2017-12-01	78	4	-6.0	36	

	Wind speed (m/s)	Visibility (10m)	Dew point temperature(°C)	\
0	2.2	2000	-17.6	
1	0.8	2000	-17.6	
2	1.0	2000	-17.7	

3	0.9	2000	-17.6
4	2.3	2000	-18.6

	Solar Radiation (MJ/m2)	Rainfall(mm)	Snowfall (cm)	Seasons	Holiday \
0	0.0	0.0	0.0	Winter	No Holiday
1	0.0	0.0	0.0	Winter	No Holiday
2	0.0	0.0	0.0	Winter	No Holiday
3	0.0	0.0	0.0	Winter	No Holiday
4	0.0	0.0	0.0	Winter	No Holiday

	Functioning Day
0	Yes
1	Yes
2	Yes
3	Yes
4	Yes

```
[50]: seoul_bike_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8760 entries, 0 to 8759
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Date                                  8760 non-null   datetime64[ns]
1   Rented Bike Count                    8760 non-null   int64
2   Hour                                 8760 non-null   int64
3   Temperature(°C)                     8760 non-null   float64
4   Humidity(%)                         8760 non-null   int64
5   Wind speed (m/s)                    8760 non-null   float64
6   Visibility (10m)                     8760 non-null   int64
7   Dew point temperature(°C)            8760 non-null   float64
8   Solar Radiation (MJ/m2)              8760 non-null   float64
9   Rainfall(mm)                        8760 non-null   float64
10  Snowfall (cm)                       8760 non-null   float64
11  Seasons                             8760 non-null   object
12  Holiday                             8760 non-null   object
13  Functioning Day                      8760 non-null   object
dtypes: datetime64[ns](1), float64(6), int64(4), object(3)
memory usage: 958.2+ KB
```

```
[51]: # Create the Day, Month, and Year quantitative attributes from the Date_
      ↪ attribute
seoul_bike_data["Day"] = seoul_bike_data["Date"].dt.day
seoul_bike_data["Month"] = seoul_bike_data["Date"].dt.month
seoul_bike_data["Year"] = seoul_bike_data["Date"].dt.year
```

```
seoul_bike_data.head()
```

```
[51]:      Date  Rented Bike Count  Hour  Temperature(°C)  Humidity(%)  \
0 2017-12-01      254      0      -5.2      37
1 2017-12-01      204      1      -5.5      38
2 2017-12-01      173      2      -6.0      39
3 2017-12-01      107      3      -6.2      40
4 2017-12-01       78      4      -6.0      36

      Wind speed (m/s)  Visibility (10m)  Dew point temperature(°C)  \
0      2.2      2000      -17.6
1      0.8      2000      -17.6
2      1.0      2000      -17.7
3      0.9      2000      -17.6
4      2.3      2000      -18.6

      Solar Radiation (MJ/m2)  Rainfall(mm)  Snowfall (cm)  Seasons  Holiday  \
0      0.0      0.0      0.0  Winter  No Holiday
1      0.0      0.0      0.0  Winter  No Holiday
2      0.0      0.0      0.0  Winter  No Holiday
3      0.0      0.0      0.0  Winter  No Holiday
4      0.0      0.0      0.0  Winter  No Holiday

      Functioning Day  Day  Month  Year
0      Yes      1      12  2017
1      Yes      1      12  2017
2      Yes      1      12  2017
3      Yes      1      12  2017
4      Yes      1      12  2017
```

```
[52]: # Drop the Date attribute as it is no longer required
seoul_bike_data.drop(columns=["Date"], inplace=True)
```

2. Seasons Attribute Second, for the Seasons attribute, I will use one hot encoding to create 4 new binary attributes to represent each Seasons value. This is because the Seasons attribute is not ordinal hence one hot encoding is a suitable method to ensure the 4 new attributes are valid for regression model building.

```
[53]: # To display all unique values of the Seasons attribute
seoul_bike_data["Seasons"].unique()
```

```
[53]: array(['Winter', 'Spring', 'Summer', 'Autumn'], dtype=object)
```

```
[54]: # Conduct the one hot encoding using pandas' get_dummies method
seoul_bike_data = pd.get_dummies(seoul_bike_data, prefix="", prefix_sep="",
    ↪columns=["Seasons"])
seoul_bike_data.head()
```

```
[54]:
```

	Rented Bike Count	Hour	Temperature(°C)	Humidity(%)	Wind speed (m/s)	\
0	254	0	-5.2	37	2.2	
1	204	1	-5.5	38	0.8	
2	173	2	-6.0	39	1.0	
3	107	3	-6.2	40	0.9	
4	78	4	-6.0	36	2.3	

	Visibility (10m)	Dew point temperature(°C)	Solar Radiation (MJ/m2)	\
0	2000	-17.6	0.0	
1	2000	-17.6	0.0	
2	2000	-17.7	0.0	
3	2000	-17.6	0.0	
4	2000	-18.6	0.0	

	Rainfall(mm)	Snowfall (cm)	Holiday	Functioning Day	Day	Month	Year	\
0	0.0	0.0	No Holiday	Yes	1	12	2017	
1	0.0	0.0	No Holiday	Yes	1	12	2017	
2	0.0	0.0	No Holiday	Yes	1	12	2017	
3	0.0	0.0	No Holiday	Yes	1	12	2017	
4	0.0	0.0	No Holiday	Yes	1	12	2017	

	Autumn	Spring	Summer	Winter
0	0	0	0	1
1	0	0	0	1
2	0	0	0	1
3	0	0	0	1
4	0	0	0	1

3. Holiday and Functioning Day Attributes Third, I will focus on transforming the Holiday and Functioning Day attributes to binary quantitative attributes. Specifically, for the Holiday attribute, if the record in our dataset is “No Holiday” then it will be mapped as 0 and if it is “Holiday” then it will be mapped as 1.

```
[55]: # To display the binary values of the Holiday attribute
seoul_bike_data["Holiday"].unique()
```

```
[55]: array(['No Holiday', 'Holiday'], dtype=object)
```

```
[56]: # Conduct the mapping using the map method and a dictionary for the binary key_
      ↪value pair
seoul_bike_data["Holiday"] = seoul_bike_data["Holiday"].map({"No Holiday":0,
      ↪"Holiday":1})
seoul_bike_data["Holiday"].unique()
```

```
[56]: array([0, 1], dtype=int64)
```

For the Functioning Day attribute, I will map the value “No” to 0 and “Yes” to 1.

```
[57]: # To display the binary values of the Functioning Day attribute
seoul_bike_data["Functioning Day"].unique()
```

```
[57]: array(['Yes', 'No'], dtype=object)
```

```
[58]: # Conduct the mapping using the map method and a dictionary for the binary key_
      ↪value pair
seoul_bike_data["Functioning Day"] = seoul_bike_data["Functioning Day"].
      ↪map({"No":0, "Yes":1})
seoul_bike_data["Functioning Day"].unique()
```

```
[58]: array([1, 0], dtype=int64)
```

4. Prepared Dataset: seoul_bike_data_prepared.csv Finally, I have completed data preparation and will now convert the dataframe containing the Seoul Bike data into a .csv flat file for the next step which is predictive modeling.

```
[59]: # To display the prepared dataset
seoul_bike_data.head()
```

```
[59]:   Rented Bike Count  Hour  Temperature(°C)  Humidity(%)  Wind speed (m/s) \
0                254    0         -5.2          37           2.2
1                204    1         -5.5          38           0.8
2                173    2         -6.0          39           1.0
3                107    3         -6.2          40           0.9
4                 78    4         -6.0          36           2.3
```

```
   Visibility (10m)  Dew point temperature(°C)  Solar Radiation (MJ/m2) \
0                2000          -17.6              0.0
1                2000          -17.6              0.0
2                2000          -17.7              0.0
3                2000          -17.6              0.0
4                2000          -18.6              0.0
```

```
   Rainfall(mm)  Snowfall (cm)  Holiday  Functioning Day  Day  Month  Year \
0             0.0           0.0        0              1    1    12  2017
1             0.0           0.0        0              1    1    12  2017
2             0.0           0.0        0              1    1    12  2017
3             0.0           0.0        0              1    1    12  2017
4             0.0           0.0        0              1    1    12  2017
```

```
   Autumn  Spring  Summer  Winter
0         0       0       0       1
1         0       0       0       1
2         0       0       0       1
3         0       0       0       1
4         0       0       0       1
```

```
[60]: # Output the prepared dataset into a .csv in the same directory as this Jupyter Notebook
seoul_bike_data.to_csv("seoul_bike_data_prepared.csv", index=False)
```

baseline_models

March 13, 2023

Predictive Modeling - Baseline Models

This Jupyter Notebook contains the 3 regression models using all attributes of the prepared dataset (seoul_bike_data_prepared.csv). Namely, linear regression, regression tree, and k-nearest neighbours. Since all attributes of the dataset will be used in the regression analyses, I have therefore named these as baseline models. 10-fold cross validation will be used.

The intent is to compare these baseline models with selected features models in terms of model performance. Selected features models in this case means regression models that I will build with only attributes that are deemed statistically significant.

Prepared Dataset

```
[2]: # Import required modules
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import cross_validate
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor
```

```
[3]: # Read in the dataset
seoul_bike = pd.read_csv("seoul_bike_data_prepared.csv", encoding="utf-8")
seoul_bike.head()
```

```
[3]:
```

	Rented Bike Count	Hour	Temperature(°C)	Humidity(%)	Wind speed (m/s)	\
0	254	0	-5.2	37	2.2	
1	204	1	-5.5	38	0.8	
2	173	2	-6.0	39	1.0	
3	107	3	-6.2	40	0.9	
4	78	4	-6.0	36	2.3	

	Visibility (10m)	Dew point temperature(°C)	Solar Radiation (MJ/m2)	\
0	2000	-17.6	0.0	
1	2000	-17.6	0.0	
2	2000	-17.7	0.0	
3	2000	-17.6	0.0	
4	2000	-18.6	0.0	

	Rainfall(mm)	Snowfall (cm)	Holiday	Functioning Day	Day	Month	Year	\
0	0.0	0.0	0	1	1	12	2017	
1	0.0	0.0	0	1	1	12	2017	
2	0.0	0.0	0	1	1	12	2017	
3	0.0	0.0	0	1	1	12	2017	
4	0.0	0.0	0	1	1	12	2017	

	Autumn	Spring	Summer	Winter
0	0	0	0	1
1	0	0	0	1
2	0	0	0	1
3	0	0	0	1
4	0	0	0	1

1. Linear Regression Model

```
[4]: # Assign the target variable, independent variables, and desired performance
      ↪ metrics
target_name = "Rented Bike Count"
target_variable = seoul_bike[target_name]
independent_variables = seoul_bike.drop(columns=target_name)
performance_metrics = ["r2", "neg_root_mean_squared_error",
      ↪ "neg_mean_absolute_error"]
```

```
[5]: # Build the linear regression model and use cross validation
lr_model = LinearRegression()
lr_scores = cross_validate(
    estimator=lr_model,
    X=independent_variables,
    y=target_variable,
    cv=10,
    scoring=performance_metrics
)
lr_scores
```

```
[5]: {'fit_time': array([0.01155519, 0.0045011 , 0.00401855, 0.0040009 , 0.00402284,
    0.00400305, 0.00452137, 0.00350046, 0.00451279, 0.00401044]),
    'score_time': array([0.00201273, 0.00149941, 0.0010004 , 0.0010004 ,
    0.00149989,
    0.00150108, 0.0009985 , 0.0010035 , 0.00149918, 0.00102329]),
    'test_r2': array([-12.37966 , -4.18163805, -0.37609285,  0.45292654,
    0.06343567, -0.01655934, -0.01606348,  0.44079453,
    0.57989526,  0.41778741]),
    'test_neg_root_mean_squared_error': array([-572.20720111, -270.87176628,
    -360.12129946, -385.88570039,
    -706.69277924, -791.0456306 , -583.69217493, -514.45991858,
    -457.16038902, -374.40065613]),
```

```
'test_neg_mean_absolute_error': array([-497.93935515, -230.05636989,
-289.2433827 , -304.93846573,
      -543.31749225, -598.69609858, -504.0723804 , -425.02535882,
      -346.68697398, -272.11677486])}]}
```

```
[6]: # Array of performance metrics scores
# Note: abs() is applied to scores returned by sklearn that are the negative
      ↳value of the metric
lr_r2_scores = lr_scores["test_r2"]
lr_root_mean_squared_error_scores =
      ↳abs(lr_scores["test_neg_root_mean_squared_error"])
lr_mean_absolute_error_scores = abs(lr_scores["test_neg_mean_absolute_error"])
```

```
[7]: # Dataframe capturing the overall performance metrics of the linear regression
      ↳model
lr_metrics = pd.DataFrame(
    {
        "Model": ["Linear Regression"],
        "R Squared": lr_r2_scores.mean(),
        "Root Mean Squared Error": lr_root_mean_squared_error_scores.mean(),
        "Mean Absolute Error": lr_mean_absolute_error_scores.mean(),
    }
)
lr_metrics
```

```
[7]:           Model  R Squared  Root Mean Squared Error  Mean Absolute Error
0  Linear Regression   -1.501517                501.653752             401.209265
```

2. Regression Tree Model

```
[8]: # Build the regression tree model and use cross validation
# random_state is set to 3 for reproducibility
rt_model = DecisionTreeRegressor(random_state=3)
rt_scores = cross_validate(
    estimator=rt_model,
    X=independent_variables,
    y=target_variable,
    cv=10,
    scoring=performance_metrics
)
rt_scores
```

```
[8]: {'fit_time': array([0.04061127, 0.0391283 , 0.04418612, 0.04155159, 0.04010677,
      0.04256129, 0.04359698, 0.03919291, 0.03850412, 0.04002833]),
      'score_time': array([0.00201488, 0.00149965, 0.00149894, 0.00150299,
      0.00150299,
      0.00150061, 0.00149894, 0.00150061, 0.00149941, 0.00149941]),
      'test_r2': array([ 0.11716893, -0.35848985, -0.01613418,  0.11199505,
```

```
0.63390278,
      0.64301151, 0.46790528, 0.58146458, 0.71809834, 0.38548109]),
      'test_neg_root_mean_squared_error': array([-146.98383877, -138.69423269,
-309.45711228, -491.63586313,
      -441.83494511, -468.77266399, -422.39417766, -445.07371382,
      -374.48861783, -384.64795771]),
      'test_neg_mean_absolute_error': array([-108.02739726, -96.21347032,
-186.07191781, -322.14383562,
      -311.36757991, -325.85388128, -279.55821918, -300.53082192,
      -229.55365297, -263.68378995])}]}
```

```
[9]: # Array of performance metrics scores
# Note: abs() is applied to scores returned by sklearn that are the negative
      ↪value of the metric
rt_r2_scores = rt_scores["test_r2"]
rt_root_mean_squared_error_scores =
      ↪abs(rt_scores["test_neg_root_mean_squared_error"])
rt_mean_absolute_error_scores = abs(rt_scores["test_neg_mean_absolute_error"])
```

```
[10]: # Dataframe capturing the overall performance metrics of the regression tree
      ↪model
rt_metrics = pd.DataFrame(
    {
        "Model": ["Regression Tree"],
        "R Squared": rt_r2_scores.mean(),
        "Root Mean Squared Error": rt_root_mean_squared_error_scores.mean(),
        "Mean Absolute Error": rt_mean_absolute_error_scores.mean(),
    }
)
rt_metrics
```

```
[10]:
```

	Model	R Squared	Root Mean Squared Error	Mean Absolute Error
0	Regression Tree	0.32844	362.398312	242.300457

3. K-Nearest Neighbours Model

```
[11]: # Build the k-nearest neighbours model and use cross validation
# I will use a range of k values (1 to 20) to determine which k contributes to
      ↪the best performing model

# Create empty lists to append each metric
knn_k = []
knn_r2 = []
knn_root_mean_squared_error = []
knn_mean_absolute_error = []

for k in range(1, 21):
    knn_model = KNeighborsRegressor(n_neighbors = k)
```

```

knn_scores = cross_validate(
    estimator=knn_model,
    X=independent_variables,
    y=target_variable,
    cv=10,
    scoring=performance_metrics
)

# Array of performance metrics scores
# Note: abs() is applied to scores returned by sklearn that are the
↳negative value of the metric
knn_r2_scores = knn_scores["test_r2"]
knn_root_mean_squared_error_scores =
↳abs(knn_scores["test_neg_root_mean_squared_error"])
knn_mean_absolute_error_scores =
↳abs(knn_scores["test_neg_mean_absolute_error"])

# Average the scores from each fold of the cross validation
# Append the metrics to lists
knn_k.append(k)
knn_r2.append(knn_r2_scores.mean())
knn_root_mean_squared_error.append(knn_root_mean_squared_error_scores.
↳mean())
knn_mean_absolute_error.append(knn_mean_absolute_error_scores.mean())

```

```

[12]: # Dataframe capturing the overall performance metrics of the k-nearest
↳neighbours model
knn_metrics = pd.DataFrame(
    {
        "Model": "K-Nearest Neighbours",
        "k": knn_k,
        "R Squared": knn_r2,
        "Root Mean Squared Error": knn_root_mean_squared_error,
        "Mean Absolute Error": knn_mean_absolute_error,
    }
)
knn_metrics

```

```

[12]:
      Model  k  R Squared  Root Mean Squared Error  \
0  K-Nearest Neighbours  1  -1.085467           596.061047
1  K-Nearest Neighbours  2  -0.586315           528.561680
2  K-Nearest Neighbours  3  -0.400623           502.123941
3  K-Nearest Neighbours  4  -0.284610           483.807584
4  K-Nearest Neighbours  5  -0.251234           475.589239
5  K-Nearest Neighbours  6  -0.210121           469.568640
6  K-Nearest Neighbours  7  -0.179231           466.370528
7  K-Nearest Neighbours  8  -0.169162           464.046049

```

8	K-Nearest Neighbours	9	-0.158443	462.093335
9	K-Nearest Neighbours	10	-0.156480	461.219524
10	K-Nearest Neighbours	11	-0.153458	460.387483
11	K-Nearest Neighbours	12	-0.149372	458.850701
12	K-Nearest Neighbours	13	-0.144513	457.627993
13	K-Nearest Neighbours	14	-0.149054	456.636305
14	K-Nearest Neighbours	15	-0.152750	456.879584
15	K-Nearest Neighbours	16	-0.154608	456.526530
16	K-Nearest Neighbours	17	-0.154704	455.835054
17	K-Nearest Neighbours	18	-0.160048	455.993327
18	K-Nearest Neighbours	19	-0.163734	456.069071
19	K-Nearest Neighbours	20	-0.166518	455.877304

	Mean Absolute Error
0	424.247146
1	386.013185
2	369.978729
3	358.960674
4	353.850411
5	350.092028
6	348.113258
7	346.585873
8	346.519673
9	346.696370
10	346.343192
11	346.010578
12	345.872111
13	345.616120
14	346.179863
15	346.261558
16	346.350020
17	346.643398
18	346.785809
19	347.192620

```
[13]: # Plot the k values by R squared to visualize the performance of each k-nearest
      ↪neighbours model
value_r2 = knn_metrics["R Squared"] == knn_metrics["R Squared"].max()
knn_metrics["colour_r2"] = np.where(value_r2 == True, "#FF7200", "#004C9B")
knn_plt = sns.regplot(
    data=knn_metrics,
    x="k",
    y="R Squared",
    fit_reg=False,
    scatter_kws={
        "alpha": 1,
        "facecolors": knn_metrics["colour_r2"],
    },
)
```

```

        "linewidths": 0,
        "s": 150,
        "zorder": 10,
    }
)

# Add title
knn_plt.set_title(
    "k-Nearest Neighbours: Optimal k Value",
    font="Arial",
    fontsize="18",
    fontweight="bold",
    loc="left"
)

# X-axis
plt.xlabel(
    "k", color="#595959", font="Arial", fontsize="14",
    ↪horizontalalignment="center"
)

# Y-axis
plt.ylabel(
    "R Squared",
    color="#595959",
    font="Arial",
    fontsize="14",
    horizontalalignment="center"
)

# Ticks
plt.xticks(range(int(knn_metrics["k"].min()), int(knn_metrics["k"].max()) + 1,
    ↪1))
plt.tick_params(colors="#595959", bottom=False, left=False, labels="14")

# Add horizontal gridlines
plt.grid(axis="y", color="#D9D9D9")

# Set plot size
knn_plt.figure.set_size_inches(14, 5)

# Spines
sns.despine(left=True)
for _, s in knn_plt.spines.items():
    s.set_color("#D9D9D9")

# Get row with max R squared

```

```

max_y_row = knn_metrics.loc[knn_metrics["R Squared"].idxmax()]


# Get the max R squared value and the corresponding x value
max_y_value = np.round(max_y_row["R Squared"], 2)
corresponding_x_value = np.round(max_y_row["k"], 2)

# Add a caption
plt.text(
    0,
    -1.35,
    f"The optimal k value for kNN is {corresponding_x_value} with a R Squared_
of {max_y_value}.",
    color="#595959",
    font="Arial",
    fontsize="14"
)

knn_plt

```

[13]: <AxesSubplot: title={'left': 'k-Nearest Neighbours: Optimal k Value'}, xlabel='k', ylabel='R Squared'>



```

[14]: # For the k-nearest neighbours models, I will choose the one with the highest R_
squared value
knn_metrics_sorted = knn_metrics.sort_values("R Squared", ascending=False)
knn_metrics_sorted.drop(columns=["k", "colour_r2"], inplace=True)

# Subset the dataframe to keep the record with the highest R Squared
highest_knn_metrics = knn_metrics_sorted.head(1)
highest_knn_metrics

```

```
[14]:
```

	Model	R Squared	Root Mean Squared Error \
12	K-Nearest Neighbours	-0.144513	457.627993

	Mean Absolute Error
12	345.872111

Performance Metrics The performance metrics of all regression models are displayed below as a pandas dataframe and a Tableau visualization.


```
[15]: # Create a baseline_metrics dataframe to capture the performance metrics of all
      ↪ regression models
baseline_metrics = pd.concat([lr_metrics, rt_metrics, highest_knn_metrics])
baseline_metrics.reset_index(drop=True, inplace=True)
baseline_metrics
```

```
[15]:
```

	Model	R Squared	Root Mean Squared Error \
0	Linear Regression	-1.501517	501.653752
1	Regression Tree	0.328440	362.398312
2	K-Nearest Neighbours	-0.144513	457.627993

	Mean Absolute Error
0	401.209265
1	242.300457
2	345.872111

```
[16]: # Export the baseline_metrics dataframe to a .csv flat file
baseline_metrics.to_csv("baseline_metrics.csv", index=False)
```



Based on the performance metrics of each model, the regression tree model is the best performing, followed by k-nearest neighbours in second place, and linear regression being last place. Notably, the R squared values for both k-nearest neighbours and linear regression are negative. As per the sklearn module, the best score for the R squared is 1.0 and a negative score means that the regression model is arbitrarily worse. In terms of the error metrics, root mean squared error and mean absolute error, the lower the error values the better performing the models are.

Therefore, it can be concluded that using all features of the prepared dataset is suboptimal for regression tree, and provides very poor performance for the k-nearest neighbours and linear regression models.

This highlights the need for variable selection using the prepared dataset (seoul_bike_data_prepared.csv) in order to optimize the performance of all regression models.

selected_features_models

March 13, 2023

Predictive Modeling - Selected Features Models

This Jupyter Notebook contains the 3 regression models using selected features of the prepared dataset (seoul_bike_data_prepared.csv). Namely, linear regression, regression tree, and k-nearest neighbours. Since select features of the dataset will be used in the regression analyses, I have therefore named these as selected features models. 10-fold cross validation will be used.

The intent is to compare these selected features models with the baseline models in terms of model performance. Selected features models in this case means regression models that I will build with only attributes that are deemed statistically significant.

Prepared Dataset

```
[1]: # Import required modules
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import cross_validate
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor
from mlxtend.feature_selection import SequentialFeatureSelector as SFS
from mlxtend.plotting import plot SequentialFeatureSelection as plot_sfs
import matplotlib.pyplot as pyplot
```

```
[2]: # Read in the dataset
seoul_bike = pd.read_csv("seoul_bike_data_prepared.csv", encoding="utf-8")
seoul_bike.head()
```

```
[2]:
```

	Rented Bike Count	Hour	Temperature(°C)	Humidity(%)	Wind speed (m/s)	\
0	254	0	-5.2	37	2.2	
1	204	1	-5.5	38	0.8	
2	173	2	-6.0	39	1.0	
3	107	3	-6.2	40	0.9	
4	78	4	-6.0	36	2.3	

	Visibility (10m)	Dew point temperature(°C)	Solar Radiation (MJ/m2)	\
0	2000	-17.6	0.0	
1	2000	-17.6	0.0	

2	2000	-17.7	0.0
3	2000	-17.6	0.0
4	2000	-18.6	0.0

	Rainfall(mm)	Snowfall (cm)	Holiday	Functioning Day	Day	Month	Year	\
0	0.0	0.0	0	1	1	12	2017	
1	0.0	0.0	0	1	1	12	2017	
2	0.0	0.0	0	1	1	12	2017	
3	0.0	0.0	0	1	1	12	2017	
4	0.0	0.0	0	1	1	12	2017	

	Autumn	Spring	Summer	Winter
0	0	0	0	1
1	0	0	0	1
2	0	0	0	1
3	0	0	0	1
4	0	0	0	1

1. Linear Regression Model

```
[3]: # Assign the target variable, independent variables, and desired performance
      ↪metrics
target_name = "Rented Bike Count"
target_variable = seoul_bike[target_name]
independent_variables = seoul_bike.drop(columns=target_name)
performance_metrics = ["r2", "neg_root_mean_squared_error",
      ↪"neg_mean_absolute_error"]
```

Feature Selection Sequential Forward Selection (SFS) will be conducted by using the mlxtend module's SequentialFeatureSelector. SFS is a sequential feature selection algorithm which automatically selects a subset of features that are the most important for the regression model. In particular, SFS adds one feature at a time based on the R squared performance metric. This reduces the model's error by discarding insignificant features and as a result, improves the performance metrics.

```
[4]: # mlxtend modules feature selection tool: Sequential Feature Selector
lr_sfs = SFS(estimator=LinearRegression(),
             k_features=(1,18),
             forward=True,
             floating=False,
             scoring="r2",
             cv=10)

# Perform the feature selection
lr_sfs.fit(independent_variables, target_variable)
```

```
[4]: SequentialFeatureSelector(cv=10, estimator=LinearRegression(),
                               k_features=(1, 18), scoring='r2')
```

```
[5]: # To visualize the Sequential Feature Selector tool results
lr_sfs_results = pd.DataFrame.from_dict(lr_sfs.get_metric_dict()).T
lr_sfs_results
```

```
[5]:
```

	feature_idx \
1	(1,)
2	(1, 4)
3	(1, 4, 7)
4	(1, 4, 7, 17)
5	(1, 4, 7, 10, 17)
6	(1, 4, 7, 10, 14, 17)
7	(1, 4, 7, 9, 10, 14, 17)
8	(1, 4, 6, 7, 9, 10, 14, 17)
9	(1, 4, 6, 7, 8, 9, 10, 14, 17)
10	(1, 4, 6, 7, 8, 9, 10, 11, 14, 17)
11	(1, 4, 6, 7, 8, 9, 10, 11, 13, 14, 17)
12	(1, 4, 6, 7, 8, 9, 10, 11, 13, 14, 15, 17)
13	(1, 4, 6, 7, 8, 9, 10, 11, 13, 14, 15, 16, 17)
14	(1, 3, 4, 6, 7, 8, 9, 10, 11, 13, 14, 15, 16, 17)
15	(1, 2, 3, 4, 6, 7, 8, 9, 10, 11, 13, 14, 15, 1...
16	(1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 13, 14, 15...
17	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 13, 14,...
18	(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13,...

	cv_scores avg_score \
1	[-0.3368353988674757, -0.918271532762825, 0.19... -0.097393
2	[-0.017461238538012402, -0.23849300655795957, ... 0.025709
3	[0.010939557944065204, -0.2462070899370501, 0... 0.040204
4	[-0.014180339709577883, -0.08067897017882486, ... 0.053138
5	[-0.03862803818220639, -0.06337005489155612, 0... 0.087882
6	[-0.0011678909621402056, -0.14506010569897887,... 0.108275
7	[0.038193597168479365, -0.1261734721658878, 0... 0.116176
8	[0.0004548344983327235, -0.13342980581946895, ... 0.115867
9	[-0.005743059846415877, -0.13461998812823261, ... 0.114459
10	[-0.025528276671164196, -0.12605749383568243, ... 0.112318
11	[-0.025528276671174632, -0.2833435904167041, 0... 0.097686
12	[-0.04250607825731101, -0.8969077226324087, 0... 0.007405
13	[-0.0425060782573079, -0.896907722632655, 0.24... 0.007405
14	[-0.6991402286332673, -1.1238234474386553, 0.1... -0.071273
15	[-1.612694820869097, -2.359029144458398, 0.174... -0.233478
16	[-1.616223784580654, -2.3300170968614142, 0.17... -0.232455
17	[-2.401920146558146, -4.164705102239599, -0.05... -0.43978
18	[-12.379659996789123, -4.181638046309049, -0.3... -1.501517

	feature_names ci_bound std_dev \
1	(Temperature(°C),) 0.242919 0.32707
2	(Temperature(°C), Visibility (10m)) 0.141344 0.190308

```

3  (Temperature(°C), Visibility (10m), Rainfall(mm))  0.152056  0.204731
4  (Temperature(°C), Visibility (10m), Rainfall(m...  0.132518  0.178424
5  (Temperature(°C), Visibility (10m), Rainfall(m...  0.174499  0.234947
6  (Temperature(°C), Visibility (10m), Rainfall(m...  0.178259  0.24001
7  (Temperature(°C), Visibility (10m), Rainfall(m...  0.18264  0.24591
8  (Temperature(°C), Visibility (10m), Solar Radi...  0.18398  0.247713
9  (Temperature(°C), Visibility (10m), Solar Radi...  0.18395  0.247673
10 (Temperature(°C), Visibility (10m), Solar Radi...  0.183854  0.247544
11 (Temperature(°C), Visibility (10m), Solar Radi...  0.198289  0.26698
12 (Temperature(°C), Visibility (10m), Solar Radi...  0.289845  0.390251
13 (Temperature(°C), Visibility (10m), Solar Radi...  0.289845  0.390251
14 (Temperature(°C), Wind speed (m/s), Visibility...  0.361759  0.487078
15 (Temperature(°C), Humidity(%), Wind speed (m/s...  0.682669  0.919157
16 (Temperature(°C), Humidity(%), Wind speed (m/s...  0.67776  0.912547
17 (Hour, Temperature(°C), Humidity(%), Wind spee...  1.106082  1.489245
18 (Hour, Temperature(°C), Humidity(%), Wind spee...  2.869342  3.863325

```

```

      std_err
1  0.109023
2  0.063436
3  0.068244
4  0.059475
5  0.078316
6  0.080003
7  0.08197
8  0.082571
9  0.082558
10 0.082515
11 0.088993
12 0.130084
13 0.130084
14 0.162359
15 0.306386
16 0.304182
17 0.496415
18 1.287775

```

```

[29]: # Global pyplot parameter
pyplot.rcParams["axes.edgecolor"] = "#D9D9D9"

# Plotting the Results
plot_lr_sfs_results = plot_sfs(lr_sfs.get_metric_dict(), kind="std_dev",
                                color="#004C9B", bcolor="#5bc2f4", figsize=[14, 5])

# Add title
pyplot.title(
    "Linear Regression Model: Sequential Forward Selection (w. StdDev)",

```

```

        font="Arial",
        fontsize="18",
        fontweight="bold",
        loc="left"
    )

    # X-axis
    pyplot.xlabel(
        "Number of Features", color="#595959", font="Arial", fontsize="14",
        ↪horizontalalignment="center"
    )

    # Y-axis
    pyplot.ylabel(
        "R Squared",
        color="#595959",
        font="Arial",
        fontsize="14",
        horizontalalignment="center"
    )

    # Ticks
    pyplot.tick_params(colors="#595959", bottom=False, left=False, labelsize="14")


    # Add horizontal gridlines
    pyplot.grid(axis="y", color="#D9D9D9")

    # Spines
    sns.despine(left=True)

    # Add a caption
    pyplot.text(
        0,
        -7.5,
        f"The highest R Squared is {np.round(lr_sfs.k_score_,3)} by using ↪
        ↪{len(lr_sfs.k_feature_idx_)} selected features.",
        color="#595959",
        font="Arial",
        fontsize="14"
    )

    pyplot.show()

```



```
[7]: # Subset the dataframe to get the record with the best R Squared results
lr_sfs_results_sorted = lr_sfs_results.sort_values("avg_score", ascending=False)
highest_lr_r2 = lr_sfs_results_sorted.head(1)
highest_lr_r2
```

```
[7]:          feature_idx \
7  (1, 4, 7, 9, 10, 14, 17)

          cv_scores avg_score \
7  [0.038193597168479365, -0.1261734721658878, 0...  0.116176

          feature_names ci_bound  std_dev \
7  (Temperature(°C), Visibility (10m), Rainfall(m...  0.18264  0.24591

          std_err
7  0.08197
```

```
[8]: # The records of the selected features
lr_selected_features =
↳ independent_variables[list(highest_lr_r2["feature_names"].values[0])]
lr_selected_features.head()
```

```
[8]:   Temperature(°C)  Visibility (10m)  Rainfall(mm)  Holiday  Functioning Day \
0             -5.2             2000             0.0         0              1
1             -5.5             2000             0.0         0              1
2             -6.0             2000             0.0         0              1
3             -6.2             2000             0.0         0              1
4             -6.0             2000             0.0         0              1

Autumn  Winter
```

```

0      0      1
1      0      1
2      0      1
3      0      1
4      0      1

```

[9]: *# Build the linear regression model and use cross validation*

```

lr_model = LinearRegression()
lr_scores = cross_validate(
    estimator=lr_model,
    X=lr_selected_features,
    y=target_variable,
    cv=10,
    scoring=performance_metrics
)
lr_scores

```

[9]: {'fit_time': array([0.00453258, 0.00300002, 0.00252867, 0.00199986, 0.00253892, 0.00251341, 0.00200129, 0.00302982, 0.00200129, 0.001508]),
'score_time': array([0.00150323, 0.00099874, 0.00150776, 0.00099993, 0.00404382, 0.00099945, 0.0010078 , 0.001508 , 0.00149894, 0.00152254]),
'test_r2': array([0.0381936 , -0.12617347, 0.29447898, 0.35963657, 0.19260796, 0.04516704, -0.45533749, 0.12238027, 0.40105209, 0.28975816]),
'test_neg_root_mean_squared_error': array([-153.41739785, -126.27949886, -257.85758852, -417.492927 , -656.1514295 , -766.6530513 , -698.56256408, -644.49422079, -545.86384631, -413.52229945]),
'test_neg_mean_absolute_error': array([-107.3295154 , -97.28586879, -201.16345441, -321.02988833, -511.21040635, -577.54951072, -619.2063464 , -554.64286717, -435.2508013 , -308.99506717])}}

[10]: *# Array of performance metrics scores*

```

# Note: abs() is applied to scores returned by sklearn that are the negative_
↳ value of the metric
lr_r2_scores = lr_scores["test_r2"]
lr_root_mean_squared_error_scores = ↳
↳ abs(lr_scores["test_neg_root_mean_squared_error"])
lr_mean_absolute_error_scores = abs(lr_scores["test_neg_mean_absolute_error"])

```

[69]: *# Dataframe capturing the overall performance metrics of the linear regression_*
↳ model

```

lr_metrics = pd.DataFrame(
    {
        "Model": ["Linear Regression"],

```

```

        "R Squared": lr_r2_scores.mean(),
        "Root Mean Squared Error": lr_root_mean_squared_error_scores.mean(),
        "Mean Absolute Error": lr_mean_absolute_error_scores.mean(),
        "Number of Features": len(lr_sfs.k_feature_idx_),
        "Feature Indices": [lr_sfs.k_feature_idx_],
        "Feature Names": [lr_sfs.k_feature_names_]
    }
)
lr_metrics

```

```

[69]:
      Model  R Squared  Root Mean Squared Error  Mean Absolute Error \
0  Linear Regression    0.116176                468.029482          373.366373

      Number of Features      Feature Indices \
0                7  (1, 4, 7, 9, 10, 14, 17)

                        Feature Names
0  (Temperature(°C), Visibility (10m), Rainfall(m...

```

2. Regression Tree Model

Feature Selection Sequential Forward Selection (SFS) will be conducted by using the `mlxtend` module's `SequentialFeatureSelector`. SFS is a sequential feature selection algorithm which automatically selects a subset of features that are the most important for the regression model. In particular, SFS adds one feature at a time based on the R squared performance metric. This reduces the model's error by discarding insignificant features and as a result, improves the performance metrics.

```

[12]: # mlxtend modules feature selection tool: Sequential Feature Selector
rt_sfs = SFS(estimator=DecisionTreeRegressor(random_state=3),
             k_features=(1,18),
             forward=True,
             floating=False,
             scoring="r2",
             cv=10)

# Perform the feature selection
rt_sfs.fit(independent_variables, target_variable)

```

```

[12]: SequentialFeatureSelector(cv=10,
                               estimator=DecisionTreeRegressor(random_state=3),
                               k_features=(1, 18), scoring='r2')

```

```

[13]: # To visualize the Sequential Feature Selector tool results
rt_sfs_results = pd.DataFrame.from_dict(rt_sfs.get_metric_dict()).T
rt_sfs_results

```


[13]:

```
feature_idx \
1      (17,)
2      (0, 17)
3      (0, 16, 17)
4      (0, 7, 16, 17)
5      (0, 7, 10, 16, 17)
6      (0, 7, 10, 13, 16, 17)
7      (0, 7, 8, 10, 13, 16, 17)
8      (0, 7, 8, 9, 10, 13, 16, 17)
9      (0, 1, 7, 8, 9, 10, 13, 16, 17)
10     (0, 1, 2, 7, 8, 9, 10, 13, 16, 17)
11     (0, 1, 2, 7, 8, 9, 10, 13, 14, 16, 17)
12     (0, 1, 2, 6, 7, 8, 9, 10, 13, 14, 16, 17)
13     (0, 1, 2, 6, 7, 8, 9, 10, 13, 14, 15, 16, 17)
14     (0, 1, 2, 5, 6, 7, 8, 9, 10, 13, 14, 15, 16, 17)
15     (0, 1, 2, 3, 5, 6, 7, 8, 9, 10, 13, 14, 15, 16,...
16     (0, 1, 2, 3, 5, 6, 7, 8, 9, 10, 11, 13, 14, 15,...
17     (0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 13, 14,...
18     (0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13,...

cv_scores avg_score \
1  [-0.06178284381364185, -0.25318375175919794, -... -0.206011
2  [0.3986806973882365, 0.19908658027429715, -0.7... 0.216475
3  [0.3986806973882365, 0.19908658027429715, -0.2... 0.299796
4  [0.4149292463114904, 0.1737705108060431, -0.45... 0.343045
5  [0.4149292463114904, 0.1737705108060431, -0.87... 0.374553
6  [0.4149292463114904, 0.11311814771476014, -0.8... 0.367048
7  [0.4075799273024534, -0.013257495179018086, -0... 0.349629
8  [0.45535361183009726, -0.3821754258535781, -0... 0.312782
9  [0.11386365245038643, -0.21352497256649006, -0... 0.301336
10 [0.1757459296888947, -0.1166019136110068, -0.1... 0.39384
11 [0.19387731539100617, -0.1520598473318615, 0.1... 0.41904
12 [0.19067333453364177, 0.01956974753667373, 0.3... 0.450626
13 [0.17654687825509563, 0.022763132631624994, 0... 0.444913
14 [0.17390822620808588, -0.060981821435387884, 0... 0.432915
15 [0.13612132012107392, -0.0557058955368015, 0.3... 0.410525
16 [0.2706576817097469, -0.09312906993473158, 0.4... 0.430233
17 [0.2903717455105228, -0.1615740565661823, 0.11... 0.393023
18 [0.11716892715250127, -0.358489850263368, -0.0... 0.32844

feature_names ci_bound std_dev \
1      (Winter,) 0.204342 0.275129
2      (Hour, Winter) 0.259982 0.350043
3      (Hour, Summer, Winter) 0.158896 0.213941
4      (Hour, Rainfall(mm), Summer, Winter) 0.215707 0.290431
5      (Hour, Rainfall(mm), Functioning Day, Summer, ... 0.326822 0.440038
6      (Hour, Rainfall(mm), Functioning Day, Year, Su... 0.33196 0.446957
```

7	(Hour, Rainfall(mm), Snowfall (cm), Functionin...	0.348485	0.469205
8	(Hour, Rainfall(mm), Snowfall (cm), Holiday, F...	0.381782	0.514037
9	(Hour, Temperature(°C), Rainfall(mm), Snowfall...	0.246263	0.331573
10	(Hour, Temperature(°C), Humidity(%), Rainfall(...	0.224895	0.302801
11	(Hour, Temperature(°C), Humidity(%), Rainfall(...	0.199242	0.268262
12	(Hour, Temperature(°C), Humidity(%), Solar Rad...	0.160101	0.215562
13	(Hour, Temperature(°C), Humidity(%), Solar Rad...	0.162683	0.219038
14	(Hour, Temperature(°C), Humidity(%), Dew point...	0.170548	0.229628
15	(Hour, Temperature(°C), Humidity(%), Wind spee...	0.171972	0.231545
16	(Hour, Temperature(°C), Humidity(%), Wind spee...	0.16262	0.218954
17	(Hour, Temperature(°C), Humidity(%), Wind spee...	0.191165	0.257387
18	(Hour, Temperature(°C), Humidity(%), Wind spee...	0.247771	0.333603

	std_err
1	0.09171
2	0.116681
3	0.071314
4	0.09681
5	0.146679
6	0.148986
7	0.156402
8	0.171346
9	0.110524
10	0.100934
11	0.089421
12	0.071854
13	0.073013
14	0.076543
15	0.077182
16	0.072985
17	0.085796
18	0.111201

```
[40]: # Global pyplot parameter
pyplot.rcParams["axes.edgecolor"] = "#D9D9D9"

# Plotting the Results
plot_rt_sfs_results = plot_sfs(rt_sfs.get_metric_dict(), kind="std_dev",
                                color="#004C9B", bcolor="#5bc2f4", figsize=[14, 5])

# Add title
pyplot.title(
    "Regression Tree Model: Sequential Forward Selection (w. StdDev)",
    font="Arial",
    fontsize="18",
    fontweight="bold",
    loc="left"
```

```

)

# X-axis
pyplot.xlabel(
    "Number of Features", color="#595959", font="Arial", fontsize="14",
    horizontalalignment="center"
)

# Y-axis
pyplot.ylabel(
    "R Squared",
    color="#595959",
    font="Arial",
    fontsize="14",
    horizontalalignment="center"
)

# Ticks
pyplot.tick_params(colors="#595959", bottom=False, left=False, labelsize="14")


# Add horizontal gridlines
pyplot.grid(axis="y", color="#D9D9D9")

# Spines
sns.despine(left=True)

# Add a caption
pyplot.text(
    0,
    -0.87,
    f"The highest R Squared is {np.round(rt_sfs.k_score_,3)} by using
    {len(rt_sfs.k_feature_idx_)} selected features.",
    color="#595959",
    font="Arial",
    fontsize="14"
)

pyplot.show()

```



```
[14]: # Subset the dataframe to get the record with the best R Squared results
rt_sfs_results_sorted = rt_sfs_results.sort_values("avg_score", ascending=False)
highest_rt_r2 = rt_sfs_results_sorted.head(1)
highest_rt_r2
```

```
[14]:          feature_idx \
12  (0, 1, 2, 6, 7, 8, 9, 10, 13, 14, 16, 17)

          cv_scores avg_score \
12  [0.19067333453364177, 0.01956974753667373, 0.3...  0.450626

          feature_names ci_bound  std_dev \
12  (Hour, Temperature(°C), Humidity(%), Solar Rad...  0.160101  0.215562

          std_err
12  0.071854
```

```
[15]: # The records of the selected features
rt_selected_features =
↳ independent_variables[list(highest_rt_r2["feature_names"].values[0])]
rt_selected_features.head()
```

```
[15]:   Hour  Temperature(°C)  Humidity(%)  Solar Radiation (MJ/m2)  Rainfall(mm) \
0      0             -5.2           37                0.0         0.0
1      1             -5.5           38                0.0         0.0
2      2             -6.0           39                0.0         0.0
3      3             -6.2           40                0.0         0.0
4      4             -6.0           36                0.0         0.0

      Snowfall (cm)  Holiday  Functioning Day  Year  Autumn  Summer  Winter
```

0	0.0	0	1	2017	0	0	1
1	0.0	0	1	2017	0	0	1
2	0.0	0	1	2017	0	0	1
3	0.0	0	1	2017	0	0	1
4	0.0	0	1	2017	0	0	1

[16]: *# Build the regression tree model and use cross validation*

```
# random_state is set to 3 for reproducibility
rt_model = DecisionTreeRegressor(random_state=3)
rt_scores = cross_validate(
    estimator=rt_model,
    X=rt_selected_features,
    y=target_variable,
    cv=10,
    scoring=performance_metrics
)
rt_scores
```

```
[16]: {'fit_time': array([0.02261305, 0.02108574, 0.02150464, 0.02118039, 0.02501607,
    0.01952267, 0.02003741, 0.02053094, 0.02002835, 0.02059603]),
    'score_time': array([0.00149989, 0.00150275, 0.00100374, 0.00150132,
    0.00149989,
    0.00252652, 0.00150156, 0.00100136, 0.0015018 , 0.00100732]),
    'test_r2': array([0.19067333, 0.01956975, 0.36915237, 0.3618115 , 0.64874401,
    0.69024584, 0.53063535, 0.57789491, 0.72323866, 0.39429686]),
    'test_neg_root_mean_squared_error': array([-140.73194994, -117.82530021,
    -243.83002257, -416.78333755,
    -432.78650869, -436.66033057, -396.71501422, -446.96769315,
    -371.0586084 , -381.87895085]),
    'test_neg_mean_absolute_error': array([-97.57534247, -78.53424658,
    -157.88584475, -281.46004566,
    -295.06164384, -299.46118721, -239.42808219, -294.96347032,
    -230.15525114, -253.41210046])}}
```

[17]: *# Array of performance metrics scores*

```
# Note: abs() is applied to scores returned by sklearn that are the negative_
↳value of the metric
rt_r2_scores = rt_scores["test_r2"]
rt_root_mean_squared_error_scores = ↳
↳abs(rt_scores["test_neg_root_mean_squared_error"])
rt_mean_absolute_error_scores = abs(rt_scores["test_neg_mean_absolute_error"])
```

[70]: *# Dataframe capturing the overall performance metrics of the regression tree_*

```
↳model
rt_metrics = pd.DataFrame(
    {
        "Model": ["Regression Tree"],
```

```

        "R Squared": rt_r2_scores.mean(),
        "Root Mean Squared Error": rt_root_mean_squared_error_scores.mean(),
        "Mean Absolute Error": rt_mean_absolute_error_scores.mean(),
        "Number of Features": len(rt_sfs.k_feature_idx_),
        "Feature Indices": [rt_sfs.k_feature_idx_],
        "Feature Names": [rt_sfs.k_feature_names_]
    }
)
rt_metrics

```

```

[70]:
      Model  R Squared  Root Mean Squared Error  Mean Absolute Error \
0  Regression Tree    0.450626                338.523772            222.793721

      Number of Features                Feature Indices \
0                12  (0, 1, 2, 6, 7, 8, 9, 10, 13, 14, 16, 17)

                        Feature Names
0  (Hour, Temperature(°C), Humidity(%), Solar Rad...

```

3. K-Nearest Neighbours Model

Feature Selection Sequential Forward Selection (SFS) will be conducted by using the mlxtend module's SequentialFeatureSelector. SFS is a sequential feature selection algorithm which automatically selects a subset of features that are the most important for the regression model. In particular, SFS adds one feature at a time based on the R squared performance metric. This reduces the model's error by discarding insignificant features and as a result, improves the performance metrics.

```

[42]: # Build the k-nearest neighbours model and use cross validation
      # I will use a range of k values (1 to 20) to determine which k contributes to
      ↪the best performing model

      # Create empty lists to append each metric
      knn_k = []
      knn_r2 = []
      knn_root_mean_squared_error = []
      knn_mean_absolute_error = []
      knn_number_of_features = []
      knn_feature_indices = []
      knn_feature_names = []

      for k in range(1, 21):

          knn = KNeighborsRegressor(n_neighbors = k)

          # mlxtend modules feature selection tool: Sequential Feature Selector
          knn_sfs = SFS(estimator=knn,

```

```

        k_features=(1,18),
        forward=True,
        floating=False,
        scoring="r2",
        cv=10)

    # Perform the feature selection
    knn_sfs.fit(independent_variables, target_variable)

    # Sequential Feature Selector tool results
    knn_sfs_results = pd.DataFrame.from_dict(knn_sfs.get_metric_dict()).T

    # Subset the dataframe to get the record with the best R Squared results
    ↪for the current k value
    knn_sfs_results_sorted = knn_sfs_results.sort_values("avg_score",
    ↪ascending=False)
    highest_knn_r2 = knn_sfs_results_sorted.head(1)

    # Capture the records of the selected features
    knn_selected_features =
    ↪independent_variables[list(highest_knn_r2["feature_names"].values[0])]

    knn_model = knn
    knn_scores = cross_validate(
        estimator=knn_model,
        X=knn_selected_features,
        y=target_variable,
        cv=10,
        scoring=performance_metrics
    )

    # Array of performance metrics scores
    # Note: abs() is applied to scores returned by sklearn that are the
    ↪negative value of the metric
    knn_r2_scores = knn_scores["test_r2"]
    knn_root_mean_squared_error_scores =
    ↪abs(knn_scores["test_neg_root_mean_squared_error"])
    knn_mean_absolute_error_scores =
    ↪abs(knn_scores["test_neg_mean_absolute_error"])
    knn_k_number_of_features = len(knn_sfs.k_feature_idx_)
    knn_k_feature_indices = knn_sfs.k_feature_idx_
    knn_k_feature_names = knn_sfs.k_feature_names_

    # Average the scores from each fold of the cross validation
    # Append the metrics to lists
    knn_k.append(k)

```

```

knn_r2.append(knn_r2_scores.mean())
knn_root_mean_squared_error.append(knn_root_mean_squared_error_scores.
↪mean())
knn_mean_absolute_error.append(knn_mean_absolute_error_scores.mean())
knn_number_of_features.append(knn_k_number_of_features)
knn_feature_indices.append(knn_k_feature_indices)
knn_feature_names.append(knn_k_feature_names)

```

```

[43]: # Dataframe capturing the overall performance metrics of the k-nearest_
↪neighbours model
knn_metrics = pd.DataFrame(
    {
        "Model": "K-Nearest Neighbours",
        "k": knn_k,
        "R Squared": knn_r2,
        "Root Mean Squared Error": knn_root_mean_squared_error,
        "Mean Absolute Error": knn_mean_absolute_error,
        "Number of Features": knn_number_of_features,
        "Feature Indices": knn_feature_indices,
        "Feature Names": knn_feature_names
    }
)
knn_metrics

```

```

[43]:

```

	Model	k	R Squared	Root Mean Squared Error \
0	K-Nearest Neighbours	1	0.219436	421.667416
1	K-Nearest Neighbours	2	0.327433	378.059800
2	K-Nearest Neighbours	3	0.381221	364.634444
3	K-Nearest Neighbours	4	0.401190	358.527755
4	K-Nearest Neighbours	5	0.413967	354.322939
5	K-Nearest Neighbours	6	0.426136	351.644301
6	K-Nearest Neighbours	7	0.430954	349.242086
7	K-Nearest Neighbours	8	0.430172	349.301746
8	K-Nearest Neighbours	9	0.428678	349.272094
9	K-Nearest Neighbours	10	0.429712	348.841335
10	K-Nearest Neighbours	11	0.430958	348.498114
11	K-Nearest Neighbours	12	0.431430	348.876643
12	K-Nearest Neighbours	13	0.428501	349.494830
13	K-Nearest Neighbours	14	0.427186	349.486821
14	K-Nearest Neighbours	15	0.425427	350.304684
15	K-Nearest Neighbours	16	0.424333	350.871335
16	K-Nearest Neighbours	17	0.421535	351.153687
17	K-Nearest Neighbours	18	0.420542	351.736602
18	K-Nearest Neighbours	19	0.420047	352.077889
19	K-Nearest Neighbours	20	0.419644	352.322494

```

Mean Absolute Error Number of Features \

```


0	275.031507	11
1	258.132192	10
2	250.500495	8
3	247.725828	7
4	245.218904	9
5	243.903063	10
6	243.780528	9
7	243.914740	8
8	244.077220	8
9	244.519863	8
10	244.424491	8
11	244.921299	8
12	245.473147	8
13	245.437565	8
14	246.307549	8
15	246.826641	8
16	247.152941	7
17	247.377194	9
18	247.762179	9
19	248.339846	9

	Feature Indices \
0	(0, 1, 6, 7, 8, 9, 10, 13, 14, 16, 17)
1	(0, 1, 5, 6, 7, 8, 10, 13, 14, 17)
2	(0, 1, 5, 7, 10, 14, 16, 17)
3	(0, 1, 5, 7, 10, 14, 17)
4	(0, 1, 5, 7, 8, 10, 14, 16, 17)
5	(0, 1, 5, 6, 7, 8, 10, 14, 16, 17)
6	(0, 1, 5, 7, 8, 10, 14, 16, 17)
7	(0, 1, 5, 7, 10, 14, 16, 17)
8	(0, 1, 5, 7, 10, 14, 16, 17)
9	(0, 1, 5, 7, 10, 14, 16, 17)
10	(0, 1, 5, 7, 10, 14, 16, 17)
11	(0, 1, 5, 7, 10, 14, 16, 17)
12	(0, 1, 5, 7, 10, 14, 16, 17)
13	(0, 1, 5, 7, 10, 14, 16, 17)
14	(0, 1, 5, 7, 10, 14, 16, 17)
15	(0, 1, 5, 7, 10, 14, 16, 17)
16	(0, 1, 5, 7, 10, 14, 17)
17	(0, 1, 5, 6, 7, 10, 14, 16, 17)
18	(0, 1, 5, 6, 7, 10, 14, 16, 17)
19	(0, 1, 5, 6, 7, 10, 14, 16, 17)

	Feature Names
0	(Hour, Temperature(°C), Solar Radiation (MJ/m2...
1	(Hour, Temperature(°C), Dew point temperature(...
2	(Hour, Temperature(°C), Dew point temperature(...

```

3  (Hour, Temperature(°C), Dew point temperature(...
4  (Hour, Temperature(°C), Dew point temperature(...
5  (Hour, Temperature(°C), Dew point temperature(...
6  (Hour, Temperature(°C), Dew point temperature(...
7  (Hour, Temperature(°C), Dew point temperature(...
8  (Hour, Temperature(°C), Dew point temperature(...
9  (Hour, Temperature(°C), Dew point temperature(...
10 (Hour, Temperature(°C), Dew point temperature(...
11 (Hour, Temperature(°C), Dew point temperature(...
12 (Hour, Temperature(°C), Dew point temperature(...
13 (Hour, Temperature(°C), Dew point temperature(...
14 (Hour, Temperature(°C), Dew point temperature(...
15 (Hour, Temperature(°C), Dew point temperature(...
16 (Hour, Temperature(°C), Dew point temperature(...
17 (Hour, Temperature(°C), Dew point temperature(...
18 (Hour, Temperature(°C), Dew point temperature(...
19 (Hour, Temperature(°C), Dew point temperature(...

```

```

[60]: # Plot the k values by R squared to visualize the performance of each k-nearest
      ↪neighbours model
value_r2 = knn_metrics["R Squared"] == knn_metrics["R Squared"].max()
knn_metrics["colour_r2"] = np.where(value_r2 == True, "#FF7200", "#004C9B")
knn_plt = sns.regplot(
    data=knn_metrics,
    x="k",
    y="R Squared",
    fit_reg=False,
    scatter_kws={
        "alpha": 1,
        "facecolors": knn_metrics["colour_r2"],
        "linewidths": 0,
        "s": 150,
        "zorder": 10,
    }
)

# Add title
knn_plt.set_title(
    "k-Nearest Neighbours: Optimal k Value",
    font="Arial",
    fontsize="18",
    fontweight="bold",
    loc="left"
)

# X-axis
plt.xlabel(

```

```

        "k", color="#595959", font="Arial", fontsize="14",
        ↪horizontalalignment="center"
    )

    # Y-axis
    plt.ylabel(
        "R Squared",
        color="#595959",
        font="Arial",
        fontsize="14",
        horizontalalignment="center"
    )

    # Ticks
    plt.xticks(range(int(knn_metrics["k"].min()), int(knn_metrics["k"].max()) + 1,
        ↪1))
    plt.tick_params(colors="#595959", bottom=False, left=False, labelsize="14")

    # Add horizontal gridlines
    plt.grid(axis="y", color="#D9D9D9")

    # Set plot size
    knn_plt.figure.set_size_inches(14, 5)

    # Spines
    sns.despine(left=True)
    for _, s in knn_plt.spines.items():
        s.set_color("#D9D9D9")

    # Get the current Axes object
    ax = plt.gca()

    # Set the ylim to begin at 0
    ax.set_ylim(top=0.5, bottom=0)

    # Get row with max R squared
    max_y_row = knn_metrics.loc[knn_metrics["R Squared"].idxmax()]

    # Get the max R squared value and the corresponding x value
    max_y_value = np.round(max_y_row["R Squared"], 2)
    corresponding_x_value = np.round(max_y_row["k"], 2)

    # Add a caption
    plt.text(
        0,
        -0.1,

```

```

    f"The optimal k value for kNN is {corresponding_x_value} with a R Squared_
of {max_y_value}.",
    color="#595959",
    font="Arial",
    fontsize="14"
)


knn_plt

```

```

[60]: <AxesSubplot: title={'left': 'k-Nearest Neighbours: Optimal k Value'},
xlabel='k', ylabel='R Squared'>

```



```

[61]: # For the k-nearest neighbours models, I will choose the one with the highest R_
squared value
knn_metrics_sorted = knn_metrics.sort_values("R Squared", ascending=False)
knn_metrics_sorted.drop(columns=["k", "colour_r2"], inplace=True)

# Subset the dataframe to keep the record with the highest R Squared
highest_knn_metrics = knn_metrics_sorted.head(1)
highest_knn_metrics

```

```

[61]:
      Model  R Squared  Root Mean Squared Error \
11  K-Nearest Neighbours    0.43143          348.876643

      Mean Absolute Error  Number of Features  Feature Indices \
11          244.921299          8  (0, 1, 5, 7, 10, 14, 16, 17)

      Feature Names
11  (Hour, Temperature(°C), Dew point temperature(...)

```

Performance Metrics The performance metrics of all regression models are displayed below as a pandas dataframe and Tableau visualizations.

```
[71]: # Create a selected_features_metrics dataframe to capture the performance
      ↪ metrics of all regression models built with selected features
selected_features_metrics = pd.concat([lr_metrics, rt_metrics,
      ↪ highest_knn_metrics])
selected_features_metrics.reset_index(drop=True, inplace=True)
selected_features_metrics
```

```
[71]:
```


	Model	R Squared	Root Mean Squared Error \
0	Linear Regression	0.116176	468.029482
1	Regression Tree	0.450626	338.523772
2	K-Nearest Neighbours	0.431430	348.876643

	Mean Absolute Error	Number of Features \
0	373.366373	7
1	222.793721	12
2	244.921299	8


	Feature Indices \
0	(1, 4, 7, 9, 10, 14, 17)
1	(0, 1, 2, 6, 7, 8, 9, 10, 13, 14, 16, 17)
2	(0, 1, 5, 7, 10, 14, 16, 17)

	Feature Names
0	(Temperature(°C), Visibility (10m), Rainfall(m...
1	(Hour, Temperature(°C), Humidity(%), Solar Rad...
2	(Hour, Temperature(°C), Dew point temperature(...

```
[74]: # Export the selected_features_metrics dataframe to a .csv flat file
selected_features_metrics.to_csv("selected_features_metrics.csv", index=False,
      ↪ encoding="utf-8-sig")
```



Based on the performance metrics of each model, the regression tree model is the best performing, followed by k-nearest neighbours in second place, and linear regression being last. Notably, the R Squared for regression tree and k-nearest neighbours are quite close in value whereas the linear regression model currently has suboptimal performance. As per the sklearn module, the best score for the R squared is 1.0. This means that the performance of the regression tree and k-nearest neighbours is approaching the halfway mark. In terms of the error metrics, root mean squared error and mean absolute error, the lower the error values the better performing the models are.



By comparing the 2 graphs, “Selected Features Models - Performance Metrics” and “Baseline Models - Performance Metrics”, we can see that using selected features in our models has improved the R squared of regression tree by approximately 0.12. In particular, for both k-nearest neighbours and linear regression, using selected features in our models has significantly improved the R squared from negative arbitrary values (very poor performance) to tangible R squared values of 0.43 and 0.12 respectively.

Therefore, by comparing the performance of the selected features models (as seen in the graph “Selected Features Models - Performance Metrics”) and the baseline models (as seen in the graph “Baseline Models - Performance Metrics”), it can be concluded that using a feature selection method, namely Sequential Forward Selection from the mlxtend module has optimized the performance of all regression models.